



CSU44061 Machine Learning

Lab 2

Adriana Hrabowych, 19304296

November 20, 2023

1 Introduction

In this assignment, the dataset used has the id '11-11-11'.

2 Part i




Figure 1: A 3D scatter plot graph visualizing part A of the lab.

2.1 (i)(a)

In (i)(a) the raw data is read in and organized into two arrays, one 2D array containing the two features and an 1D array containing the y values. Then the training data is visualized on the 3d graph seen in figure one using blue circular markers. Each data point is placed depending on the value of its two features

and the target value, with the X-axis corresponding to X_1 , the Y-axis corresponding to X_2 and the Z-axis corresponding to the target y value.

When graphed, the training data takes on the shape of a curved plane.

2.2 (i)(b)

As seen in section (i)(a) our data lies on a curve and as such is non linear, so in order for our models to be more accurate we need to train them with additional polynomial features. In this lab we added polynomial features up to a factor of five by using sklearn's PolynomialFeatures function. This gives us a set of 21 features made up of variations on X_1 and X_2 which are:

$$\theta_{(1-21)} = 1, X_1, X_2, X_1^2, X_1X_2, X_2^2, X_1^3, X_1^2X_2, X_1X_2^2, X_2^3, X_1^4, X_1^3X_2, X_1^2X_2^2, X_1X_2^3, X_2^4, X_1^5, X_1^4X_2, X_1^3X_2^2, X_1^2X_2^3, X_1X_2^4, X_2^5$$

Now using this new set of polynomial features, four Lasso regression models were training. In Lasso, the models are linear with the addition of a cost function. The model itself is $h_\theta(x) = \theta^T x$. The objective of the model is to minimize the cost function:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_\theta(x^{(i)}), y^{(i)}) + \frac{\lambda}{m} \sum_{j=1}^n |\theta_j|$$

Where C is a weight hyperparameter. In sklearn α is used instead of C , and $\alpha = 1/(2C)$. The only difference between the four models trained is that they each have different C 's: 1, 10, 100, and 1000. After training, the feature parameters were:

Lasso	θ_1	θ_2	θ_3	θ_4	θ_5	θ_6	θ_7	θ_8	θ_9	θ_{10}	θ_{11}
$C = 1$	0	0	0	0	0	0	0	0	0	0	0
$C = 10$	0	0	-0.8517	0.5652	0	0	0	0	0	0	0
$C = 100$	0	-0.03534	-0.9898	1.0261	0	0	0	0	0	0	0
$C = 1,000$	0	-0.0087	-0.9829	1.0018	-0.0581	0	-0.0641	-0.0743	0	0.0257	0.0547

Figure 2: A table containing the Lasso feature parameters values for θ_1 to θ_{11} in the columns with each row representing a model with a different C value.

Lasso	θ_{12}	θ_{13}	θ_{14}	θ_{15}	θ_{16}	θ_{17}	θ_{18}	θ_{19}	θ_{20}	θ_{21}	Intercept
$C = 1$	0	0	0	0	0	0	0	0	0	0	0.3981
$C = 10$	0	0	0	0	0	0	0	0	0	0	0.1712
$C = 100$	0	0	0	0	0	0	0	0	0	0	0.0067
$C = 1,000$	0	0.074	0.1676	0.0376	0	0	0	-0.0418	0.028	0	-0.0138

Figure 3: A table containing the Lasso feature parameters values for θ_{12} to θ_{21} and the intercept in the columns with each row representing a model with a different C value.

As seen in figures 2 and 3, generally as the C is raised the absolute value of each parameter also increases as the features have more influence on the predictions. The opposite is true for the intercept, which grows smaller as C increases. A lot of the feature parameters are zero as Lasso models automatically perform feature selection to eliminate the weights of less important features.

2.3 (i)(c)

Instead of splitting our data into training and testing sections, we trained our models on all of the raw data and predict using a grid of feature values. As the original values for X_1 and X_2 were between -1 and 1, I chose to use a grid from [-3,3] as to not extend too far past the original values while still being able to see beyond the original data. The predictions are then plotted on a 3D plane.

As seen in figure 4, as C is increased the predictions begin to change from a flat linear plane (as seen when $C=1$) to a curved surface (seen in all other C values graphed). We originally stated that the data laid on


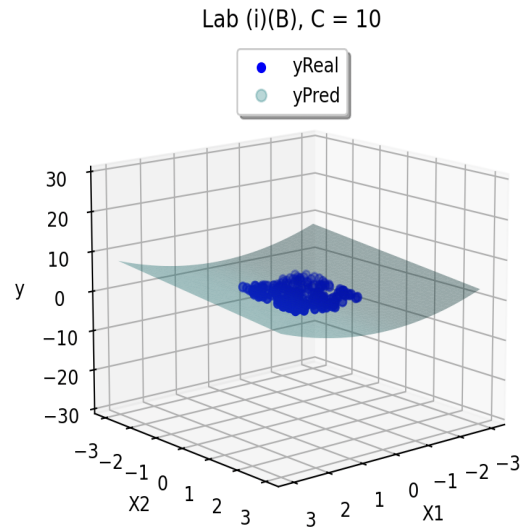
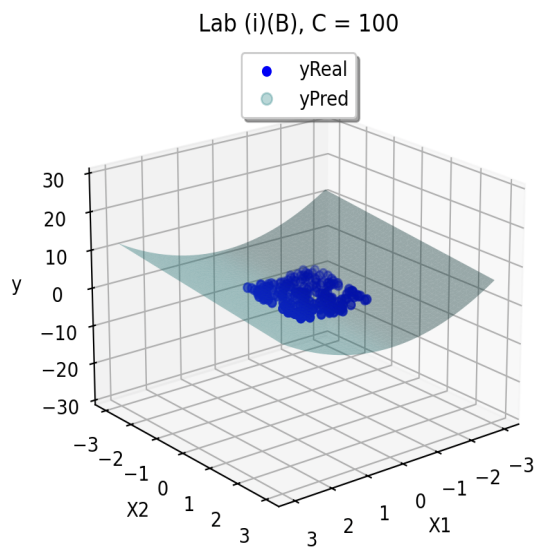

(a) Lasso model predictions with $C = 1$ (b) Lasso model predictions with $C = 10$ (c) Lasso model predictions with $C = 100$ (d) Lasso model predictions with $C = 1,000$

Figure 4: A figure with four subfigures, each containing a 3D scatter plot and plane that visualizes part ib of the lab.

a curved plane, so the models with higher C's seem to more accurately portray this. Though as the C gets much larger, the simple curve gives way to more complex structures.

2.4 (i)(d)

Underfitting is when too few features are used and as such the model cannot accurately portray the behavior of the original data. Underfit models are often incredibly inaccurate because of this. Overfitting is when too many features are used and the 'noise' in the data starts to have more influence on the predictions the model makes and as such it cannot categorize data correct.

A higher C value means the model will place a higher weight on the training data, which means that too high of a C value will cause the model to be overfit as it becomes 'overly accurate'. A lower C value puts less weight on the training data and allows more inaccuracies in predictions as it generalizes. Too low of a C can cause the lasso model to eliminate the weights of every feature parameters to 0, which causes it to be underfit. This can be seen in the graphs of figure 4, as when C=1 the surface is flat and linear and too simple to fit the data well. Whereas when C=1,000, the predicted surface becomes more complex and bumpy as it fits the data too well.

2.5 (i)(e)

Once again using the same set of polynomial features we now train four Ridge Regression models. Ridge models are linear with the addition of a penalty cost function. The objective for Ridge models is to minimize a residual sum of squares function:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

Where C is a complexity parameter, represented by α in sklearn, which also equals $1/2C$. The only difference between the four models trained is that they each have different C's: 1, 0.1, 0.01, 0.001. After training, the feature parameters were:

Ridge	θ_1	θ_2	θ_3	θ_4	θ_5	θ_6	θ_7	θ_8	θ_9	θ_{10}	θ_{11}
C = 1	0	-0.0078	-0.9526	0.7613	-0.0925	-0.16	-0.1143	-0.1363	-0.0324	0.0409	0.2834
C = 0.1	0	-0.0205	-0.7368	0.532	0.0065	0.0033	-0.0404	-0.1558	0.0083	-0.1914	0.3971
C = 0.01	0	-0.0285	-0.4236	0.2354	0.01	0.0061	-0.034	-0.1263	-0.0062	-0.2221	0.2002
C = 0.001	0	-0.0122	-0.1121	0.0362	0.0048	0.0032	-0.0096	-0.0362	-0.0036	-0.0676	0.0307

Figure 5: A table containing the Ridge feature parameters values for θ_1 to θ_{11} in the columns with each row representing a model with a different C value.

Ridge	θ_{12}	θ_{13}	θ_{14}	θ_{15}	θ_{16}	θ_{17}	θ_{18}	θ_{19}	θ_{20}	θ_{21}	Intercept
C = 1	0.0505	0.1742	0.1914	0.0347	0.0636	0.0463	0.0145	-0.0427	0.088	-0.0512	0.015
C = 0.1	-0.0074	0.1686	0.059	0.0012	-0.0187	-0.0576	-0.0089	-0.0267	0.0195	-0.0663	0.0702
C = 0.01	-0.0057	0.0799	0.0158	0.0015	-0.0287	-0.0707	-0.0121	-0.0608	-0.0024	-0.1479	0.2323
C = 0.001	0.0012	0.0126	0.005	0.0023	-0.0069	-0.0214	-0.0029	-0.0204	-0.0017	-0.0486	0.3672

Figure 6: A table containing the Ridge feature parameters values for θ_{12} to θ_{21} and the intercept in the columns with each row representing a model with a different C value.

In the Ridge models as C gets bigger the feature parameters grow smaller. Unlike Lasso, as C gets larger the intercepts also get larger. The main difference between the two is the Ridge models have less parameters equal to 0 as Ridge Regression does not automatically eliminate the weight of features of lesser importance. Again the models are tested using the same grid of X values from $[-3,3]$ and plotted using a 3D surface, as seen in figure 7. The C value in ridge models controls the strength of the l2 penalty, as C increases the variance decreases and vice versa. The C also is in charge of keeping parameters small.

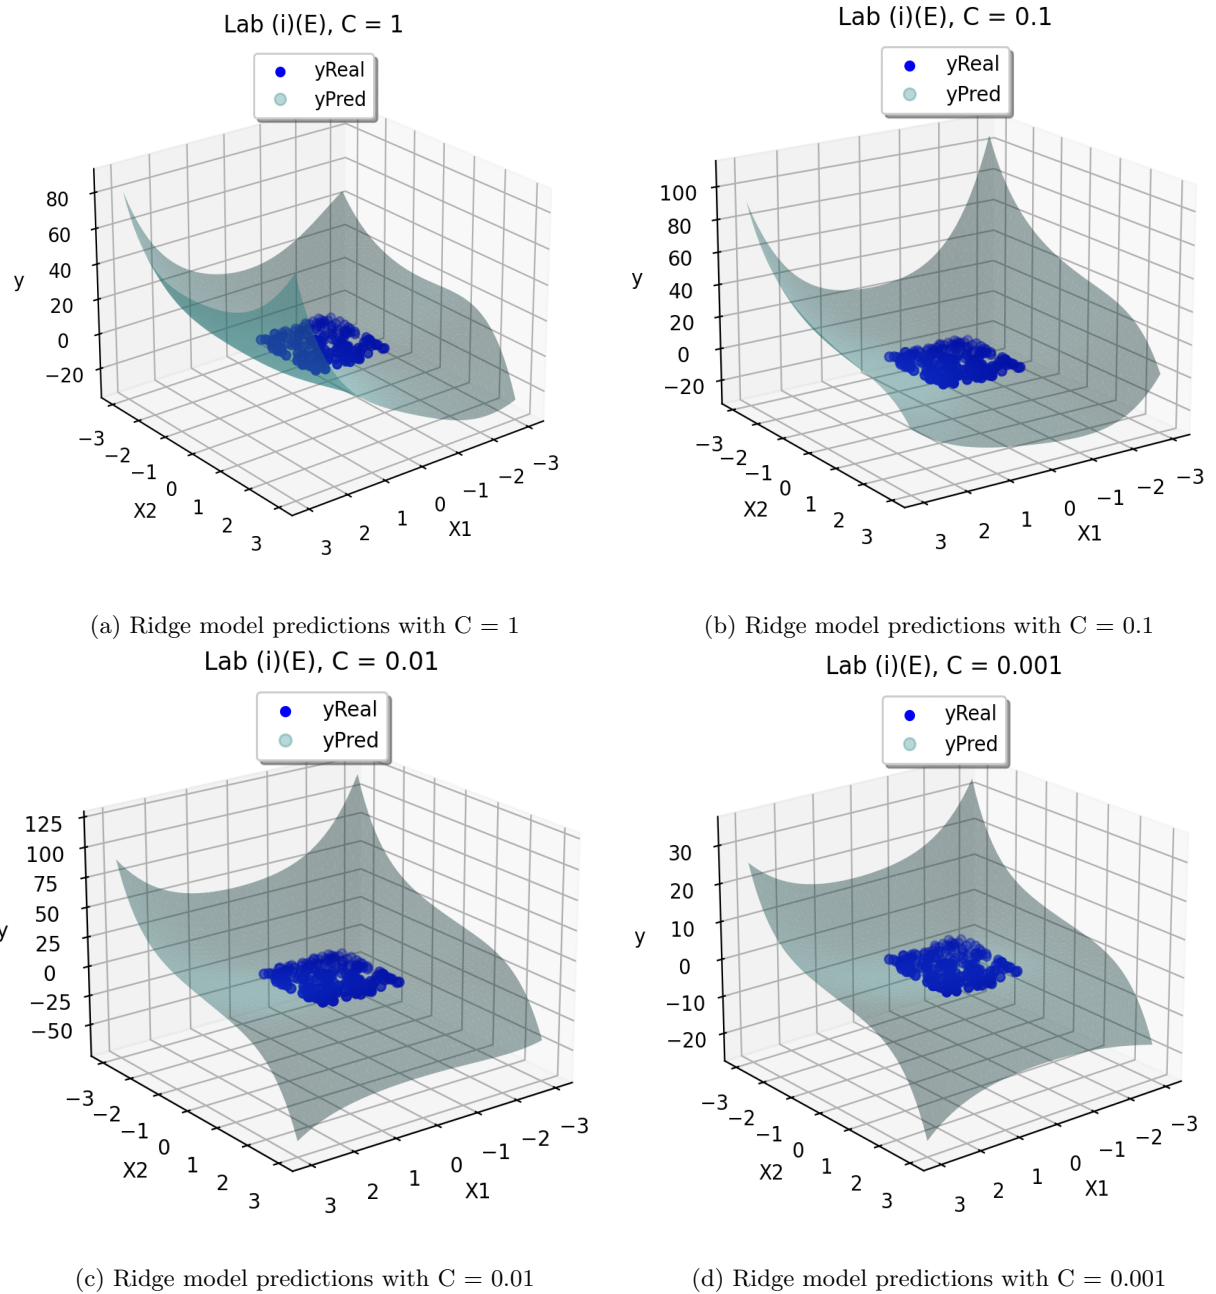


Figure 7: A figure with four subfigures, each containing a 3D scatter plot and plane that visualizes part ie of the lab.

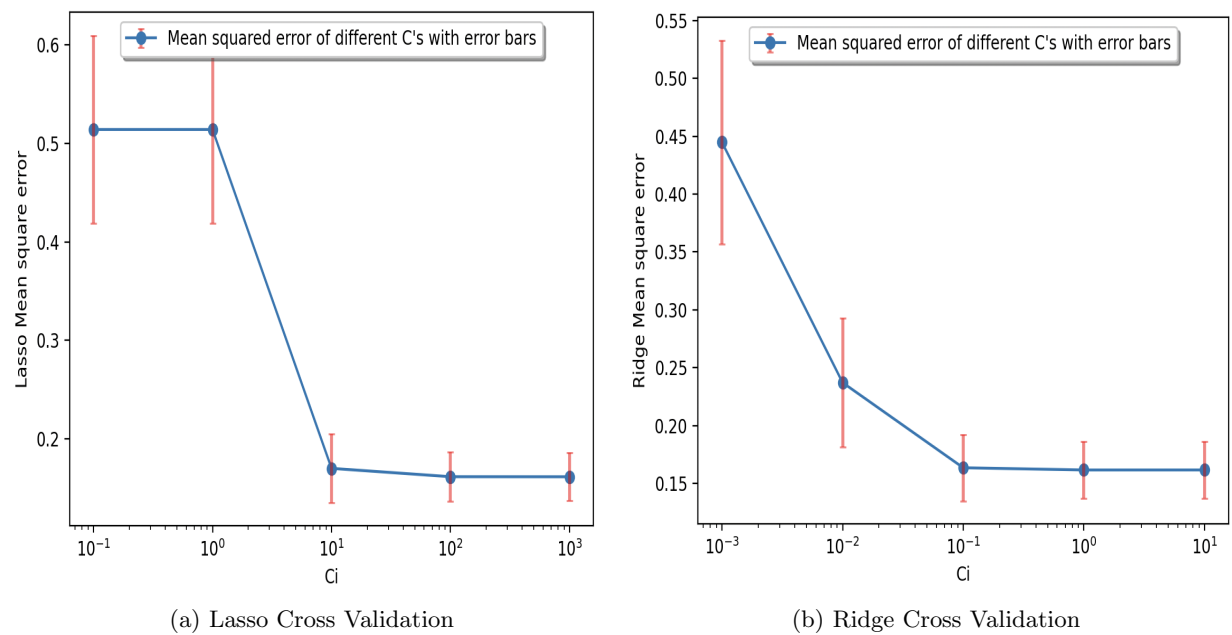


Figure 8: A figure with two subfigures, each containing a mean error bar plot that visualizes part iia of the lab.

3 Part ii

3.1 (ii)(a)

We now perform K-fold Cross Validation on Lasso models in order to see what value of C is the best one to choose. We are using a 5-fold as specified in the instructions, so we split the data into five equal parts and train the model on four of these parts and then test on the final part. We repeat this until each part has had the chance to be the testing data. With each iteration we take the mean squared error of the predictions and once the iterations are done we calculated the mean of these mean squared errors and the standard deviation.

Then we repeat this for new models with different C values, the C values I chose for this cross validation are $[0.1, 1, 10, 100, 1000]$. I chose these values as they are the C values I used for my models in previous parts, plus 0.1 which I added to explore going even smaller. Also, these C 's extend small enough and large enough to explore both underfitting and overfitting.

3.2 (ii)(b)

Once we have the mean mean squared error and the standard deviation of these errors for each value of C , we then plot these means as points on a graph with the standard deviation creating an error bar. In figure 8.a we see the results of our Lasso Cross Validation. We can see as C gets larger, our mean and standard deviation gets smaller. It's important to note that as C increases, α in sklearn decreases. From these results it looks like a C of 10 would be the best, as any lower and our mean error is too high as we are underfit while any higher will probably overfit our data.

3.3 (ii)(c)

We then repeated the K-fold Cross Validation with Ridge Regression models and with different values of C . The values I chose for Ridge were $[0.001, 0.01, 0.1, 1, 10]$. I chose these values as they are the same as the ones used in my models from part ie with the addition of 10 to explore larger C 's. Also because these represent a wide spread of C values that are both common and uncommon for Ridge models.

We see our Ridge cross validation results in figure 8.b. From this graph we can see that the best value of C

for our Ridge model is 0.01 or 0.1, as any larger will probably overfit our data and any smaller will underfit as seen by the large mean errors and error bars.

4 Appendix

```

1 # id:11-11--11
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from sklearn.linear_model import Lasso, Ridge
6 from sklearn.metrics import mean_squared_error
7 from sklearn.preprocessing import PolynomialFeatures
8 from sklearn.model_selection import KFold
9
10 ##Function that adds the raw data to a graph
11 def addRawData(graph, name, X, Y):
12     graph.scatter(X[:,0], X[:,1], Y, color='b', marker='o')
13     graph.set_title(name)
14     graph.set_xlabel("X1")
15     graph.set_ylabel("X2")
16     graph.set_zlabel("y")
17     if name != "Lab (i)(a)":
18         graph.plot([0],[0], linestyle="none", c='teal', alpha=0.3, marker = 'o')
19         graph.legend(["yReal","yPred"], loc='upper center',fancybox=True, shadow=True)
20     else:
21         graph.legend(["yReal"], loc='upper center',fancybox=True, shadow=True)
22
23
24
25 # (i)(a)
26 ##Read in data
27 df = pd.read_csv("week3.csv")
28 X1=df.iloc[:,0]
29 X2=df.iloc[:,1]
30 X=np.column_stack((X1,X2))
31 y=df.iloc[:,2]
32
33 ##Graph Data Points
34 fig = plt.figure()
35 ax = fig.add_subplot(111, projection='3d')
36 addRawData(ax, "Lab (i)(a)", X, y)
37 plt.show()
38
39 # (i)(b)
40 ##Add polynomial features
41 polyX = PolynomialFeatures(5).fit_transform(np.array(X))
42
43 ##Train Model
44 lasso1 = Lasso(alpha=0.5) #C = 1
45 lasso2 = Lasso(alpha=0.05) #C = 10
46 lasso3 = Lasso(alpha=0.005) #C = 100
47 lasso4 = Lasso(alpha=0.0005) #C = 1000
48
49 lasso1.fit(polyX, y)
50 lasso2.fit(polyX, y)
51 lasso3.fit(polyX, y)
52 lasso4.fit(polyX, y)
53
54 ##Print Coefficients and Intercepts
55 print("Lasso C=1 Coefficients")
56 for i in range(21):
57     print("Coefficient theta-", i+1, ": ", lasso1.coef_[i])
58
59 print("Lasso C=10 Coefficients")
60 for i in range(21):
61     print("Coefficient theta-", i+1, ": ", lasso2.coef_[i])
62

```

```

63 print("Lasso C=100 Coefficients")
64 for i in range(21):
65     print("Coefficient theta-", i+1, ": ", lasso3.coef_[i])
66
67 print("Lasso C=1,000 Coefficients")
68 for i in range(21):
69     print("Coefficient theta-", i+1, ": ", lasso4.coef_[i])
70
71 print("Lasso C=1 Intercept")
72 print(lasso1.intercept_)
73 print("Lasso C=100 Intercept")
74 print(lasso2.intercept_)
75 print("Lasso C=1,000 Intercept")
76 print(lasso3.intercept_)
77 print("Lasso C=10,000 Intercept")
78 print(lasso4.intercept_)
79
80 #(i)(c)
81
82 xGrid = []
83 grid = np.linspace(-3,3)
84 for i in grid:
85     for j in grid:
86         xGrid.append([i,j])
87 xGrid = np.array(xGrid)
88 xGrid = PolynomialFeatures(5).fit_transform(xGrid)
89
90 ##Predict
91 ypred1 = lasso1.predict(xGrid)
92 ypred2 = lasso2.predict(xGrid)
93 ypred3 = lasso3.predict(xGrid)
94 ypred4 = lasso4.predict(xGrid)
95
96 ##Create Graphs
97 fig = plt.figure()
98 ax = fig.add_subplot(111, projection='3d')
99 addRawData(ax, "Lab (i)(B), C = 1", X, y)
100 ax.plot_trisurf(xGrid[:,1], xGrid[:,2], ypred1, alpha=0.4, color='teal')
101 ax.set_zlim([-30, 30])
102 plt.show()
103
104 fig = plt.figure()
105 ax = fig.add_subplot(111, projection='3d')
106 addRawData(ax, "Lab (i)(B), C = 10", X, y)
107 ax.plot_trisurf(xGrid[:,1], xGrid[:,2], ypred2, alpha=0.4, color='teal')
108 ax.set_zlim([-30, 30])
109 plt.show()
110
111 fig = plt.figure()
112 ax = fig.add_subplot(111, projection='3d')
113 addRawData(ax, "Lab (i)(B), C = 100", X, y)
114 ax.plot_trisurf(xGrid[:,1], xGrid[:,2], ypred3, alpha=0.4, color='teal')
115 ax.set_zlim([-30, 30])
116 plt.show()
117
118 fig = plt.figure()
119 ax = fig.add_subplot(111, projection='3d')
120 addRawData(ax, "Lab (i)(B), C = 1,000", X, y)
121 ax.plot_trisurf(xGrid[:,1], xGrid[:,2], ypred4, alpha=0.4, color='teal')
122 ax.set_zlim([-30, 30])
123 plt.show()
124
125 #(i)(e)
126 ##Create and fit models
127 ridge1 = Ridge(alpha=0.5) #C = 1
128 ridge2 = Ridge(alpha=5.0) #C = 0.1
129 ridge3 = Ridge(alpha=50.0) #C = 0.01
130 ridge4 = Ridge(alpha=500.0) #C = 0.001
131

```



```

132 ridge1.fit(polyX, y)
133 ridge2.fit(polyX, y)
134 ridge3.fit(polyX, y)
135 ridge4.fit(polyX, y)
136
137 ##Print Coefficients and Intercepts
138 print("Ridge C=1 Coefficients")
139 for i in range(21):
140     print("Coefficient theta-", i+1, ": ", ridge1.coef_[i])
141
142 print("Ridge C=0.1 Coefficients")
143 for i in range(21):
144     print("Coefficient theta-", i+1, ": ", ridge2.coef_[i])
145
146 print("Ridge C=0.01 Coefficients")
147 for i in range(21):
148     print("Coefficient theta-", i+1, ": ", ridge3.coef_[i])
149
150 print("Ridge C=0.001 Coefficients")
151 for i in range(21):
152     print("Coefficient theta-", i+1, ": ", ridge4.coef_[i])
153
154 print("Ridge C=1 Intercept")
155 print(ridge1.intercept_)
156 print("Ridge C=0.1 Intercept")
157 print(ridge2.intercept_)
158 print("Ridge C=0.01 Intercept")
159 print(ridge3.intercept_)
160 print("Ridge C=0.001 Intercept")
161 print(ridge4.intercept_)
162
163 ##Predict
164 ypred1 = ridge1.predict(xGrid)
165 ypred2 = ridge2.predict(xGrid)
166 ypred3 = ridge3.predict(xGrid)
167 ypred4 = ridge4.predict(xGrid)
168
169 ##Create Graphs
170 fig = plt.figure()
171 ax = fig.add_subplot(111, projection='3d')
172 addRawData(ax, "Lab (i)(E), C = 1", X, y)
173 ax.plot_trisurf(xGrid[:,1], xGrid[:,2], ypred1, alpha=0.4, color='teal')
174 plt.show()
175
176 fig = plt.figure()
177 ax = fig.add_subplot(111, projection='3d')
178 addRawData(ax, "Lab (i)(E), C = 0.1", X, y)
179 ax.plot_trisurf(xGrid[:,1], xGrid[:,2], ypred2, alpha=0.4, color='teal')
180 plt.show()
181
182 fig = plt.figure()
183 ax = fig.add_subplot(111, projection='3d')
184 addRawData(ax, "Lab (i)(E), C = 0.01", X, y)
185 ax.plot_trisurf(xGrid[:,1], xGrid[:,2], ypred3, alpha=0.4, color='teal')
186 plt.show()
187
188 fig = plt.figure()
189 ax = fig.add_subplot(111, projection='3d')
190 addRawData(ax, "Lab (i)(E), C = 0.001", X, y)
191 ax.plot_trisurf(xGrid[:,1], xGrid[:,2], ypred4, alpha=0.4, color='teal')
192 plt.show()
193
194 #(ii)(a)
195
196 mean_error=[]; std_error=[]
197 Ci_range = [0.1, 1, 10, 100, 1000]
198 for Ci in Ci_range:
199     model = Lasso(alpha=1/(2*Ci))
200     temp=[]

```

```
201     kf = KFold(n_splits=5)
202     for train, test in kf.split(polyX):
203         model.fit(X[train], y[train])
204         ypred = model.predict(X[test])
205         temp.append(mean_squared_error(y[test], ypred))
206     mean_error.append(np.array(temp).mean())
207     std_error.append(np.array(temp).std())
208
209 markers, caps, bars = plt.errorbar(Ci_range, mean_error, barsabove=True, yerr=std_error,
210     ecolor='r', capsize=2, elinewidth=2, fmt="-o")
211 [bar.set_alpha(0.5) for bar in bars]
212 [cap.set_alpha(0.5) for cap in caps]
213
214 plt.xlabel('Ci')
215 plt.legend(["Mean squared error of different C's with error bars"], loc='upper center',
216     fancybox=True, shadow=True)
217 plt.ylabel('Lasso Mean square error')
218 plt.xscale('log')
219 plt.show()
220
221 # (ii) (c)
222
223 mean_error=[]; std_error=[]
224 Ci_range = [0.001, 0.01, 0.1, 1, 10]
225 for Ci in Ci_range:
226     model = Ridge(alpha=1/(2*Ci))
227     temp=[]
228     kf = KFold(n_splits=5)
229     for train, test in kf.split(polyX):
230         model.fit(polyX[train], y[train])
231         ypred = model.predict(polyX[test])
232         temp.append(mean_squared_error(y[test], ypred))
233     mean_error.append(np.array(temp).mean())
234     std_error.append(np.array(temp).std())
235
236 markers, caps, bars = plt.errorbar(Ci_range, mean_error, barsabove=True, yerr=std_error,
237     ecolor='r', capsize=2, elinewidth=2, fmt="-o")
238 [bar.set_alpha(0.5) for bar in bars]
239 [cap.set_alpha(0.5) for cap in caps]
240
241 plt.xlabel('Ci')
242 plt.legend(["Mean squared error of different C's with error bars"], loc='upper center',
243     fancybox=True, shadow=True)
244 plt.ylabel('Ridge Mean square error')
245 plt.xscale('log')
246 plt.show()
```